

PG外部表应用

mysql_fdw, oss_fdw

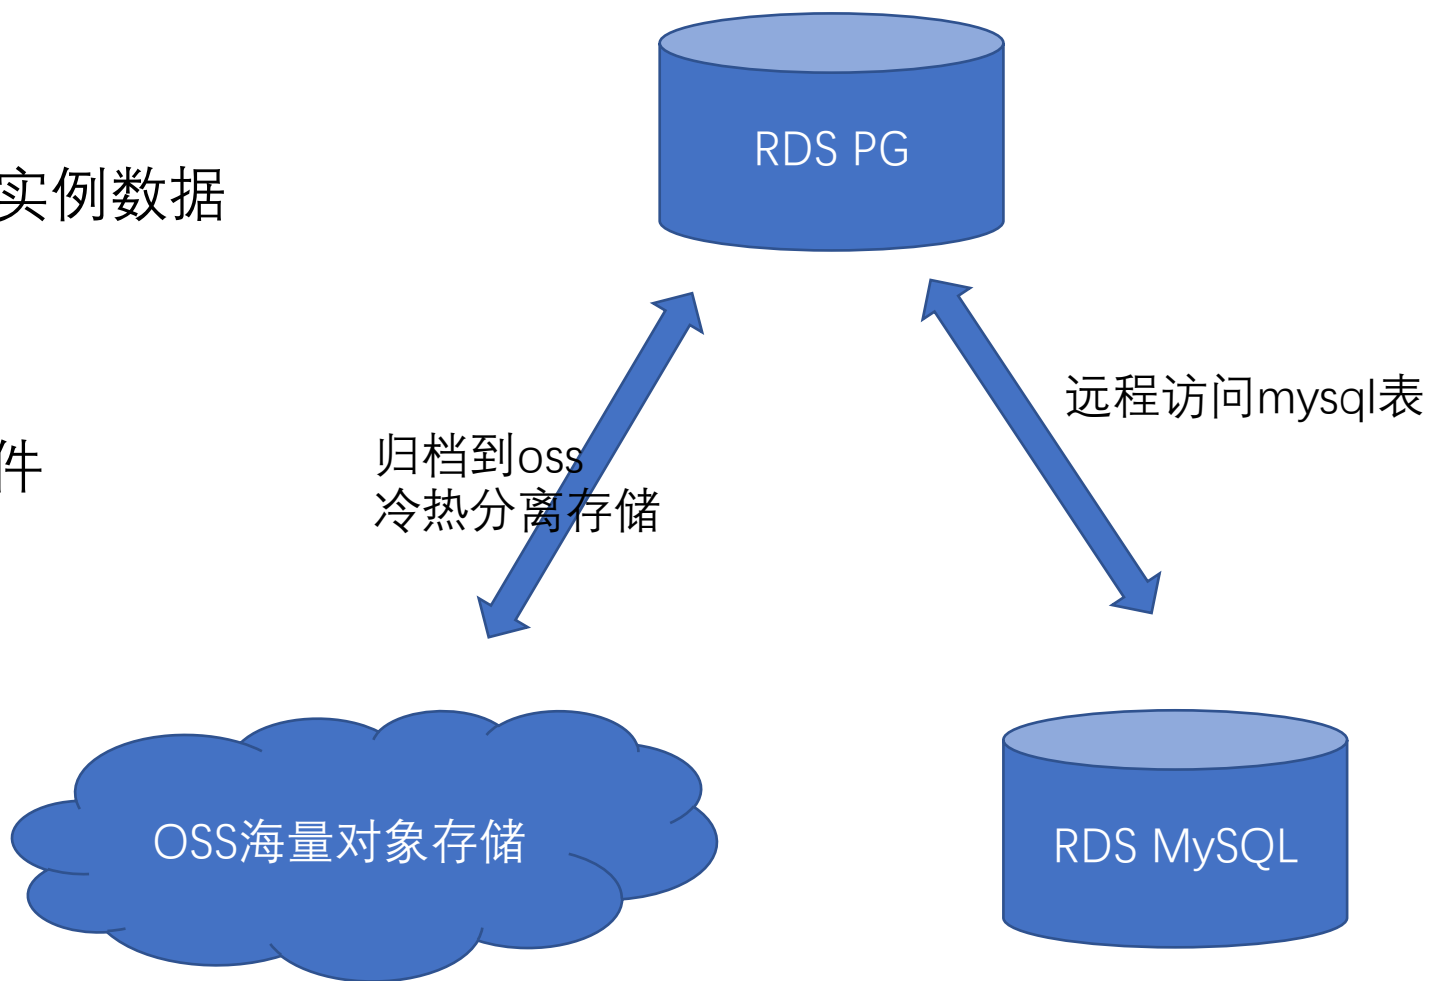
(直接读写MySQL数据, 冷热分离、外部归档表)

阿里云

digoal

目录

- mysql_fdw
 - 从外部表直接读写mysql实例数据
- oss_fdw
 - 冷热分离
 - 归档历史数据到oss
 - 从外部表直接读写oss文件



mysql_fdw

- server(host:port)
 - mysql的网络配置
- user mapping(user:pwd)
 - pg数据库user如何连接到以上mysql server (mysql的用户密码)
- foreign table(dbname,tablename,column define)
 - mysql表结构如何映射到pg外部表

环境

- RDS MySQL
- RDS PG 12
 - 相同vpc , vswitch
- pg实例可以通过vpc内网访问mysql server
- mysql数据, 使用第5课的测试表和数据

使用第5课生成数据

- test_mm
- test_innodb

创建插件、mysql server、映射用户密码

使用rds_superuser创建 extension

```
create extension mysql_fdw ;
```

创建 mysql server

```
CREATE SERVER mysql_1
```

```
    FOREIGN DATA WRAPPER mysql_fdw
```

```
    OPTIONS (host 'mysql网络地址', port 'mysql端口');
```

设置 pg数据库用户digoal的user mapping

```
CREATE USER MAPPING FOR digoal
```

```
    SERVER mysql_1
```

```
    OPTIONS (username 'mysqluser', password 'mysqluserpwd');
```

方法1， 单独建外部表

```
CREATE FOREIGN TABLE t1(  
id INT8 NOT NULL,  
user_id VARCHAR (20) NOT NULL,  
group_id INT8 NOT NULL,  
create_time timestamp NOT NULL  
)  
SERVER mysql_1  
  OPTIONS (dbname 'db1', table_name 'test_mm');
```

```
db1=> select count(*) from t1;
```

```
count
```

```
-----
```

```
1000000
```

```
(1 row)
```

方法2，一次性import所有外部表

```
db1=> create schema ft;
```

```
CREATE SCHEMA
```

```
db1=> import foreign schema db1 from server mysql_1 into ft;
```

```
IMPORT FOREIGN SCHEMA
```

```
db1=> \det ft.*
```

List of foreign tables

Schema	Table	Server
--------	-------	--------

-----+-----+-----		
-------------------	--	--

ft	test_innodb	mysql_1
----	-------------	---------

ft	test_mm	mysql_1
----	---------	---------

(2 rows)

```
db1=> select count(*) from ft.test_mm;
```

count

1000000

(1 row)

查询、观察远程sql

```
db1=> select * from t1 where id=2;
```

id	user_id	group_id	create_time
2	qGZKzux5FVrfN5RUN4QO	19	2020-01-08 15:07:21

(1 row)

```
db1=> explain verbose select * from t1 where id=2;
```

QUERY PLAN

Foreign Scan on digoal.t1 (cost=25.00..1025.00 rows=1000 width=82)

Output: id, user_id, group_id, create_time

Remote server startup cost: 25

Remote query: SELECT `id`, `user_id`, `group_id`, `create_time` FROM `db1`.`test_mm` **WHERE ((`id` = 2))**

(4 rows)

条件没有下推，会导致传输所有记录

db1=> explain verbose select * from t1 limit 100;

QUERY PLAN

Limit (cost=25.00..125.00 rows=100 width=82)

Output: id, user_id, group_id, create_time

-> Foreign Scan on digoal.t1 (cost=25.00..1025.00 rows=1000 width=82)

Output: id, user_id, group_id, create_time

Remote server startup cost: 25

Remote query: SELECT `id`, `user_id`, `group_id`, `create_time` FROM `db1`.`test_mm`

(6 rows)

插入远程mysql

```
db1=> explain verbose insert into t1 (user_id, group_id, create_time) values ('digoal', 123, now());
```

QUERY PLAN

Insert on digoal.t1 (cost=0.00..0.01 rows=1 width=82)

-> Result (cost=0.00..0.01 rows=1 width=82)

Output: NULL::bigint, 'digoal'::character varying(20), '123'::bigint, now()

(3 rows)

```
db1=> insert into t1 (user_id, group_id, create_time) values ('digoal', 123, now());
```

INSERT 0 1

```
db1=> select * from t1 where user_id='digoal';
```

id	user_id	group_id	create_time
----	---------	----------	-------------

1000003	digoal	123	2020-01-08 16:11:05
---------	--------	-----	---------------------

(1 row)

更新远程mysql

```
db1=> explain verbose update t1 set group_id=321 where id=1000003;
```

QUERY PLAN

Update on digoal.t1 (cost=25.00..1025.00 rows=1000 width=90)

-> Foreign Scan on digoal.t1 (cost=25.00..1025.00 rows=1000 width=90)

Output: id, user_id, '321'::bigint, create_time, id

Remote server startup cost: 25

Remote query: SELECT `id`, `user_id`, `create_time` FROM `db1`.`test_mm` WHERE ((`id` = 1000003)) FOR UPDATE

(5 rows)

```
db1=> update t1 set group_id=321 where id=1000003;
```

UPDATE 1

```
db1=> select * from t1 where user_id='digoal';
```

id | user_id | group_id | create_time

-----+-----+-----+-----

1000003 | digoal | 321 | 2020-01-08 16:11:05

(1 row)

删除远程mysql记录

```
db1=> explain verbose delete from t1 where id=1000003;
```

QUERY PLAN

Delete on digoal.t1 (cost=25.00..1025.00 rows=1000 width=8)

-> Foreign Scan on digoal.t1 (cost=25.00..1025.00 rows=1000 width=8)

Output: id

Remote server startup cost: 25

Remote query: SELECT `id` FROM `db1`.`test_mm` WHERE ((`id` = 1000003)) FOR UPDATE

(5 rows)

```
db1=> delete from t1 where id=1000003;
```

DELETE 1

```
db1=> select * from t1 where user_id='digoal';
```

id | user_id | group_id | create_time

----+-----+-----+-----

(0 rows)

注意

- 检查explain verbose, 如果没有push down, 需要返回全表数据到pg

查询有哪些外部表、server、用户映射

```
db1=> select * from pg_foreign_server ;
```

oid	srvname	srvowner	srvfdw	srvtype	srvversion	srvacl	srvoptions
16561	mysql_1	16385	16559				{host=rm-xxxx.mysql.rds.aliyuncs.com,port=3306}

(1 row)

```
db1=> \h alter server
```

Command: **ALTER SERVER**

Description: change the definition of a foreign server

Syntax:

```
ALTER SERVER name [ VERSION 'new_version' ]
```

```
[ OPTIONS ( [ ADD | SET | DROP ] option ['value'] [, ... ] ) ]
```

```
ALTER SERVER name OWNER TO { new_owner | CURRENT_USER | SESSION_USER }
```

```
ALTER SERVER name RENAME TO new_name
```

URL: <https://www.postgresql.org/docs/12/sql-alterserver.html>

查询有哪些外部表、server、用户映射

```
db1=> \des
```

List of foreign servers

Name	Owner	Foreign-data wrapper
------	-------	----------------------

mysql_1	digoal	mysql_fdw
ossserver1	digoal	oss_fdw

(2 rows)

```
db1=> \deu
```

List of user mappings

Server	User name
--------	-----------

mysql_1	digoal
---------	--------

(1 row)

```
db1=> \det
```

List of foreign tables

Schema	Table	Server
--------	-------	--------

public	oss_tb1	ossserver1
public	oss_tb2	ossserver1
public	test_mm	mysql_1

(3 rows)

归档数据、冷热分离

- oss_fdw
 - 目前仅支持PG V10, 未来会覆盖所有主流版本 (11,12等)

申请有oss访问权限的AK



申请有oss访问权限的AK



快速创建子用户AccessKey



STEP 1: 填写用户名

STEP 2: 选择权限

STEP 3: 创建用户及AK

* 用户名:

pg_oss_fdw_user

长度1-64个字符，允许输入大小写英文字母、数字、"."、 "_"或"-"

下一步

快速创建子用户AccessKey



STEP 1: 填写用户名

STEP 2: 选择权限

STEP 3: 创建用户及AK

oss|

系统 AliyunOSSFullAccess

管理对象存储服务(OSS)权限

系统 AliyunOSSReadOnlyAccess

只读访问对象存储服务(OSS)的权限

系统 AliyunYundunNewBGPAntiDDoSserv...

管理云盾新BGP高防IP(New BGP Ant...

系统 AliyunYundunNewBGPAntiDDoSserv...

只读访问新BGP高防IP(New BGP Ant...

上一步

开始创建

快速创建子用户AccessKey



STEP 1: 填写用户名

STEP 2: 选择权限

STEP 3: 创建用户及AK

这是AccessKey可供下载的唯一机会，请及时保存！



用户 pg_oss_fdw_user 创建成功

您已成功新建用户和AccessKey，并且添加完成授权。

AccessKey详情



AccessKeyId:

[Redacted AccessKeyId]

AccessKeySecret:

[Redacted AccessKeySecret]

查看用户详情

下载AccessKey

申请oss bucket

对象存储 / 概览

存储空间 + ↻ ↕

Bucket 名称

基础数据

仅统计当前 Bucket 合计数据，平均延迟 1~2 小时。不作为计量数据，仅作参考。不同存储类型参考帮助文档

存储用量

总用量 (不含 ECS 快照)

76.59 TB

本月流量

外网流出流量

458.23 GB

Endpoint

oss-cn-hangzhou.aliyuncs.com

创建 Bucket

注意：Bucket 创建成功后，您所选择的 存储类型、区域 不支持变更。

Bucket 名称

pgarchive01

区域

华东1 (杭州)

Endpoint

oss-cn-hangzhou.aliyuncs.com

存储类型

标准存储

低频访问

归档存储

读写权限

私有

公共读

公共读写

服务器端加密

无

AES256

KMS

确定

取消

<https://oss.console.aliyun.com/>

服务器端加密

无

AES256

KMS

i 将文件上传至 OSS 后，自动对其进行落盘加密存储，KMS 加密方式需要进行权限设置，当前 KMS 仅支持 OSS 默认托管的 CMK，如需试用 KMS 单独创建的 CMK 加密（BYOK），请和我们联系，了解 [更多服务器端加密指南](#)。

实时日志查询

开通

不开通

OSS 与日志服务深度结合，免费提供最近 7 天内的 OSS 实时日志查询。开通该功能后，您可对 Bucket 的访问记录进行实时查询分析，[了解详情](#)。

确定

取消

对象存储

概览

存储空间

Bucket 名称

0 Byte

月同比 -- 日环比 --

0 Byte

上月外网流出流量: 0 Byte

0

上月请求次数: 0

0

0

基础数据

基础数据统计平均延迟 1-2 小时。不作为计量数据，仅作参考。

存储用量

总用量 (不含 ECS 快照)

本月流量

外网流出流量

本月请求次数

读请求

文件数量

文件碎片

访问域名

	EndPoint (地域节点)	Bucket 域名	HTTPS
外网访问	oss-cn-hangzhou.aliyuncs.com	pgarchive01.oss-cn-hangzhou.aliyuncs.com	支持
ECS 的经典网络访问 (内网)	oss-cn-hangzhou-internal.aliyuncs.com	pgarchive01.oss-cn-hangzhou-internal.aliyuncs.com	支持
ECS 的 VPC 网络访问 (内网)	oss-cn-hangzhou-internal.aliyuncs.com	pgarchive01.oss-cn-hangzhou-internal.aliyuncs.com	支持
传输加速域名 (全地域上传下载加速)	未开启	<div>开启</div>	支持

基础设置

oss_fdw 采用ecs的vpc网络访问（内网） EndPoint

创建oss_fdw, 创建server

PostgreSQL 创建插件

```
create extension oss_fdw;
```

创建 server

```
CREATE SERVER ossserver1 FOREIGN DATA WRAPPER oss_fdw OPTIONS
```

```
(host 'oss-cn-hangzhou-internal.aliyuncs.com', id 'accessid内容', key 'accesskey内容', bucket  
'pgarchive01');
```

创建oss外部归档表

创建 oss 外部表，压缩

```
CREATE FOREIGN TABLE oss_tb1
(date text, time text, open float,
high float, low float, volume int)
SERVER ossserver1
OPTIONS ( dir 'oss_tb1/', delimiter ',',
format 'csv', encoding 'utf8', compressiontype 'gzip' );
```

不压缩

```
CREATE FOREIGN TABLE oss_tb2
(date text, time text, open float,
high float, low float, volume int)
SERVER ossserver1
OPTIONS ( dir 'oss_tb2/', delimiter ',',
format 'csv', encoding 'utf8', compressiontype 'none' );
```

创建本地表， 写入测试数据

创建表， 数据就装载到这张表中

```
create table tbl1
```

```
(date text, time text, open float,  
    high float, low float, volume int) ;
```

```
insert into tbl1 select md5(random()::text), now()::text,  
random()*10000, random()*100000000, random()*10000,  
random()*1000000000  
from generate_series(1,1000000);
```

将本地表数据写入外部表

数据从 ossexample 装载到 example 中。

```
db1=> insert into oss_tb1 select * from tbl1;
```

WARNING: oss compress process does not close

WARNING: oss compres thread does not close

NOTICE: begin writiing data to oss location oss_tb1/, with block size 10 MB and oss file size 1024 MB

INSERT 0 1000000

Time: 8899.910 ms (00:08.900)

查询外部表

```
db1=> select * from oss_tb1 limit 1;
```

NOTICE: a total of 1 files will be loaded, begin oss_tb1/_oss_tb1_631787927722928, end oss_tb1/_oss_tb1_631787927722928, file format gzip

date	time	open	high	low	volume
ee421482662f0424dfb337925b35f029	2020-01-08 16:35:07.795619+08	4588.77929486334	3359252.90361047	3289.49342481792	47414201

(1 row)

Time: 198.820 ms

```
db1=> select count(*) from oss_tb1;
```

NOTICE: a total of 1 files will be loaded, begin oss_tb1/_oss_tb1_631787927722928, end oss_tb1/_oss_tb1_631787927722928, file format gzip

count
1000000

(1 row)

Time: 5873.628 ms (00:05.874)

查询外部表

```
db1=> insert into oss_tb1 select * from tbl1;
```

```
NOTICE: begin writiing data to oss location oss_tb1/, with block size 10 MB and oss file size 1024 MB
```

```
INSERT 0 1000000
```

```
Time: 8720.626 ms (00:08.721)
```

```
db1=> select count(*) from oss_tb1;
```

```
NOTICE: a total of 2 files will be loaded, begin oss_tb1/_oss_tb1_631787927722928, end  
oss_tb1/_oss_tb1_631788054925016, file format gzip
```

```
count
```

```
-----
```

```
2000000
```

```
(1 row)
```

```
Time: 12469.130 ms (00:12.469)
```

查询外部表

```
db1=> insert into oss_tb2 select * from tbl1;
```

```
NOTICE: begin writiing data to oss location oss_tb2/, with block size 10 MB and oss file size 1024 MB
```

```
INSERT 0 1000000
```

```
Time: 7120.737 ms (00:07.121)
```

```
db1=> insert into oss_tb2 select * from tbl1;
```

```
NOTICE: begin writiing data to oss location oss_tb2/, with block size 10 MB and oss file size 1024 MB
```

```
INSERT 0 1000000
```

```
Time: 7004.067 ms (00:07.004)
```

```
db1=> select count(*) from oss_tb2;
```

```
NOTICE: a total of 2 files will be loaded, begin oss_tb2/_oss_tb2_631788259206276, end oss_tb2/_oss_tb2_631788267353423, file format text
```

```
count
```

```
-----
```

```
2000000
```

```
(1 row)
```

```
Time: 21546.740 ms (00:21.547)
```


压缩与不压缩的存储空间区别

```
db1=> select * from oss_fdw_list_file('oss_tb1','public');
```

name	size
------	------

-----+-----	
-------------	--

oss_tb1/_oss_tb1_631787927722928	53922166
----------------------------------	----------

oss_tb1/_oss_tb1_631788054925016	53922166
----------------------------------	----------

(2 rows)

```
db1=> select * from oss_fdw_list_file('oss_tb2','public');
```

name	size
------	------

-----+-----	
-------------	--

oss_tb2/_oss_tb2_631788259206276	122555480
----------------------------------	-----------

oss_tb2/_oss_tb2_631788267353423	122555480
----------------------------------	-----------

(2 rows)

归档数据不支持记录级别的更新和删除

```
db1=> update oss_tb1 set date=now() where volume=1;
```

```
ERROR: cannot update foreign table "oss_tb1"
```

```
db1=> delete from oss_tb1 where volume=1;
```

```
ERROR: cannot delete from foreign table "oss_tb1"
```

```
db1=> delete from oss_tb2 where volume=1;
```

```
ERROR: cannot delete from foreign table "oss_tb2"
```

```
db1=> update oss_tb2 set date=now() where volume=1;
```

```
ERROR: cannot update foreign table "oss_tb2"
```

扩展资料

- 阿里云RDS PostgreSQL OSS 外部表实践 - (dblink异步调用封装并行) 从OSS并行导入数据
- https://github.com/digoal/blog/blob/master/201804/20180427_01.md
- 阿里云RDS PostgreSQL OSS 外部表实践 - (dblink异步调用封装并行) 数据并行导出到OSS
- https://github.com/digoal/blog/blob/master/201709/20170906_01.md
- 强制数据分布与导出prefix - 阿里云pg, hdb pg oss快速数据规整外部表导出实践案例
- https://github.com/digoal/blog/blob/master/201801/20180109_01.md

注意事项

- https://help.aliyun.com/document_detail/44461.html
 - filepath, dir, prefix 参数互斥
 - 有oss写入要求时, 必须使用dir或prefix
 - 有oss写入要求时, 不能使用parse_errors

参考资料

- mysql_fdw手册
 - https://github.com/EnterpriseDB/mysql_fdw
- oss_fdw手册
 - https://help.aliyun.com/document_detail/44461.html
- MySQL手册
 - <https://www.mysqltutorial.org/>
 - <https://dev.mysql.com/doc/refman/8.0/en/>
- PG 管理、开发规范
 - https://github.com/digoal/blog/blob/master/201609/20160926_01.md
- PG手册
 - <https://www.postgresql.org/docs/current/index.html>
 - <https://www.postgresqltutorial.com/postgresql-tutorial/postgresql-vs-mysql/>
- GIS手册
 - <http://postgis.net/docs/manual-3.0/>

一期开课计划(PG+MySQL联合方案)

- - 2019.12.30 19:30 RDS PG产品概览，如何与MySQL结合使用
- - 2019.12.31 19:30 如何连接PG， GUI， CLI的使用
- - 2020.1.3 19:30 如何压测PG数据库、如何瞬间构造海量测试数据
- - 2020.1.6 19:30 MySQL与PG对比学习(面向开发者)
- - 2020.1.7 19:30 如何将MySQL数据同步到PG (DTS)
- - 2020.1.8 19:30 PG外部表妙用 - mysql_fdw, oss_fdw (直接读写MySQL数据、冷热分离)
- - 2020.1.9 19:30 PG应用场景介绍 - 并行计算， 实时分析
- - 2020.1.10 19:30 PG应用场景介绍 - GIS
- - 2020.1.13 19:30 PG应用场景介绍 - 用户画像、实时营销系统
- - 2020.1.14 19:30 PG应用场景介绍 - 多维搜索
- - 2020.1.15 19:30 PG应用场景介绍 - 向量计算、图像搜索
- - 2020.1.16 19:30 PG应用场景介绍 - 全文检索、模糊查询
- - 2020.1.17 19:30 PG 数据分析语法介绍
- - 2020.1.18 19:30 PG 更多功能了解：扩展语法、索引、类型、存储过程与函数。如何加入PG技术社群

本课程习题

- 如何一次性创建所有mysql的表作为pg的外部表
- 可以删除mysql外部表的数据吗
- 如何查看访问外部表的远程SQL
- 访问频次非常非常低的归档数据，使用什么方法存储更加廉价
- oss外部归档数据支持记录级别的删除和更新吗
- 如何提升归档数据写入oss的速度
- 如何提升查询的oss归档表速度

技术社群



PG技术交流钉钉群(3500+人)

